1 point

1. (True/False) k-means always converges to a local optimum.

- (●) True
- ( ) False

---

1 point

2. (True/False) The clustering objective is non-increasing throughout a run of k-means.

- (●) True
- ( ) False

---

1 point

3. (True/False) Running k-means with a larger value of k always enables a lower possible final objective value than running k-means with smaller k.

- (●) True
- ( ) False

---

1 point

4. (True/False) Any initialization of the centroids in k-means is just as good as any other.

- ( ) True
- (●) False

---

1 point

5. (True/False) Initializing centroids using k-means++ guarantees convergence to a global optimum.

- ( ) True
- (●) False

---

1 point

6. (True/False) Initializing centroids using k-means++ costs more than random initialization in the beginning, but can pay off eventually by speeding up convergence.

- (●) True
- ( ) False

---

1 point

7. (True/False) Using k-means++ can only influence the number of iterations to convergence, not the quality of the final assignments (i.e., objective value at convergence).

- ( ) True
- (●) False

---

4 points

8. Consider the following dataset:

	X1	X2
Data point 1	-1.88	2.05
Data point 2	-0.71	0.42
Data point 3	2.41	-0.67
Data point 4	1.85	-3.80
Data point 5	-3.69	-1.33

Perform k-means with k=2 until the cluster assignment does not change between successive iterations. Use the following initialization for the centroids:

	X1	X2
Cluster 1	2.00	2.00
Cluster 2	-2.00	-2.00

Which of the five data points changed its cluster assignment most often during the k-means run?

- ( ) Data point 1
- (●) Data point 2
- ( ) Data point 3
- ( ) Data point 4
- ( ) Data point 5

---

1 point

9. Suppose we initialize k-means with the following centroids

Which of the following best describes the cluster assignment in the first iteration of k-means?

(●)

( )

( )

( )

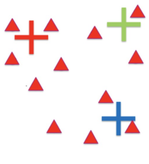

( )